

# Assignment Project Exam Help

Databases: Introduction

<https://tutorcs.com>

P.J. McBrien

Imperial College London

WeChat: cstutorcs

# Databases are Computer Stores of Data!

Tiny Bank Ltd Customer: McBrien, P.  
Strand Branch Current Acc: 10000100  
Sortcode: 55-66-67

Trans	Amount	Date
1000	2300.00	5/1/1999
1002	1233.41	3/1/1999
1006	10.23	15/1/1999

Tiny Bank Ltd Customer: Poulovassilis, A.  
Wimbledon Branch Current Acc: 10000107  
Sortcode: 55-66-56

Trans	Amount	Date
1004	-116.00	11/1/1999
1007	315.56	15/1/1999

Tiny Bank Ltd Customer: McBrien, P.  
Strand Branch Deposit Acc: 10000101  
Sortcode: 55-66-67

Trans	Amount	Date
1001	4000.00	5/1/1999
1008	1230.00	15/1/1999

Tiny Bank Ltd Customer: Poulovassilis, A.  
Wimbledon Branch Deposit Acc: 10000119  
Sortcode: 55-66-56

Trans	Amount	Date
1009	5600.00	18/1/1999

Tiny Bank Ltd Customer: Boyd, M.  
Goodge St Branch Current Acc: 10000103  
Sortcode: 55-66-34

Trans	Amount	Date
1005	145.50	12/1/1999

Tiny Bank Ltd Customer: Bailey, J.  
Wimbledon Branch Current Acc: 10000125  
Sortcode: 55-66-56

Trans	Amount	Date
No transactions this month		

Deposit Rates

AccountRate

101 5.25

119 5.50

# Relational Data Model

## Relational Data Model

Roughly: storing data in tables

bank_data									
no	sortcode	lname	cash	type	sname	rate?	<u>mid</u>	amount	tdate
100	67	Strand	34005.00	current	McBrien, P.		1000	2300.00	1999-01-05
101	67	Strand	34005.00	deposit	McBrien, P.	5.25	1001	4000.00	1999-01-05
100	67	Strand	34005.00	current	McBrien, P.		1002	-223.45	1999-01-08
107	56	Wimbledon	84340.45	current	Poulovassilis, A.		1004	-100.00	1999-01-11
103	34	Geodge St	6900.67	current	Boyd, M.		1005	145.50	1999-01-12
100	67	Strand	34005.00	current	McBrien, P.		1006	10.23	1999-01-15
107	56	Wimbledon	84340.45	current	Poulovassilis, A.		1007	345.56	1999-01-15
101	67	Strand	34005.00	deposit	McBrien, P.	5.25	1008	1230.00	1999-01-15
119	56	Wimbledon	84340.45	deposit	Poulovassilis, A.	5.50	1009	5600.00	1999-01-18

## Database Design: ER Modelling

# Assignment Project Exam Help

<https://tutorcs.com>

WeChat: tutorcs



## Structured Data: Relational Model

branch		
sortcode	branch	cash
56	'Wimbledon'	94340.45
34	'Goodge St'	8900.67
67	'Strand'	34005.00

movement			
mid	no	amount	tdate
1000	100	2300.00	5/1/1999
1001	101	4000.00	5/1/1999
1002	100	-223.45	8/1/1999
1004	107	-100.00	11/1/1999
1005	103	145.50	12/1/1999
1006	100	10.23	15/1/1999
1007	107	345.56	15/1/1999
1008	101	1230.00	15/1/1999
1009	119	5600.00	18/1/1999

account				
no	type	aname	rate?	sortcode
100	'current'	'McBrien, P.'	NULL	67
101	'deposit'	'McBrien, P.'	5.25	67
103	'current'	'Boyd, M.'	NULL	34
107	'current'	'Poulovassilis, A.'	NULL	56
119	'deposit'	'Poulovassilis, A.'	5.50	56
125	'current'	'Bailey, J.'	NULL	56

key branch(sortcode)

key branch(bname)

key movement(mid)

key account(no)

movement(no)  $\xRightarrow{fk}$  account(no)account(sortcode)  $\xRightarrow{fk}$  branch(sortcode)

## Data Model: CSV

## Assignment Project Exam Help

branch.csv  
 sortcode,bname,cash  
 56,"Wimbledon",94340.45  
 34,"Goodge St", 8900.67  
 67,"Strand",34005.00

account.csv

no,type,cname,rate,sortcode  
 100,"current","McBrien, P.",,67  
 101,"deposit","McBrien, P.",5.25,67  
 103,"current","Boyd, M.",,34  
 107,"current","Poulovassilis, A.",,56  
 119,"deposit","Poulovassilis, A.",5.50,56  
 125,"current","Bailey, J.",,56

movement.csv  
 mid,no,amount,tdate  
 1000,100,2300.00,5/1/1999  
 1001,101,4000.00,5/1/1999  
 1002,100,-223.45,8/1/1999  
 1004,107, 100.00,11/1/1999  
 1005,103,145.50,12/1/1999  
 1006,100,10.23,15/1/1999  
 1007,107,345.56,15/1/1999  
 1008,101,1230.00,15/1/1999  
 1009,119,5600.00,18/1/1999

<https://tutorcs.com>  
 WeChat: cstutorcs

## Semistructured Data: XML

# Assignment Project Exam Help

```
<bank>
  <branch sortcode="67" bname="Strand" cash="34005.00">
    <account no="100" type="current" cname="McBrien, P.">
      <movement mid="1000" amount="2300.00" tdate="5/1/1999" />
      <movement mid="1002" amount="223.45" tdate="8/1/1999" />
      <movement mid="1005" amount="15.23" tdate="15/1/1999" />
    </account>
    <account no="101" type="deposit" cname="McBrien, P." rate="5.25">
      <movement mid="1001" amount="4000.00" tdate="5/1/1999" />
      <movement mid="1008" amount="1230.00" tdate="15/1/1999" />
    </account>
  </branch>
</bank>
```

<https://tutorcs.com>

WeChat: cstutorcs

# SQL DDL: Implementation of the Relational Model

```
CREATE TABLE branch
( sortcode INTEGER NOT NULL,
  bname VARCHAR(20) NOT NULL,
  cash DECIMAL(10,2) NOT NULL,
  CONSTRAINT branch_pk PRIMARY KEY (sortcode)
)
```

```
CREATE UNIQUE INDEX branch_bname_idx
ON branch(bname)
```

```
CREATE TABLE account
( no INTEGER NOT NULL,
  type CHAR(8) NOT NULL,
  cname VARCHAR(20) NOT NULL,
  rate DECIMAL(4,2) NOT NULL,
  sortcode INTEGER NOT NULL,
  CONSTRAINT account_pk
    PRIMARY KEY (no),
  CONSTRAINT account_fk
    FOREIGN KEY (sortcode) REFERENCES branch
)

CREATE INDEX account_type_idx ON account(type)
```

```
CREATE TABLE movement
( mid INTEGER NOT NULL,
  no INTEGER NOT NULL,
  amount DECIMAL(10,2) NOT NULL,
  tdate DATETIME NOT NULL,
  CONSTRAINT movement_pk
    PRIMARY KEY (mid),
  CONSTRAINT movement_fk
    FOREIGN KEY (no) REFERENCES account
)
```



## SQL DML: Implementation of the Relational Algebra

## Basic SQL SELECT statements

```
SELECT no, cname, rate
FROM account
WHERE type='deposit'
```

Assignment Project Exam Help

## SQL Joins

```
SELECT bname, no, rate
FROM branch JOIN account USING (sortcode)
WHERE type='deposit'
```

Same as

```
SELECT bname, no, rate
FROM account JOIN branch ON branch.sortcode=account.sortcode
WHERE type='deposit'
```

Same as

```
SELECT bname, no, rate
FROM account, branch
WHERE branch.sortcode=account.sortcode
AND type='deposit'
```

## RDBMS Products

## Assignment Project Exam Help

Product	SQL Language	Company
DB2	SQL PL	IBM
Oracle	PL/SQL	Oracle
Sybase	Transact-SQL	SAP
SQLServer	Transact-SQL	Microsoft
PostgreSQL	PL/pgSQL	Open Source
MySQL	MySQL	Open Source (Oracle)

<https://tutorcs.com>

WeChat: cstutorcs

*All partially implement ANSI SQL*

# Transactions

```

BEGIN TRANSACTION
UPDATE branch
SET cash = cash - 10000.00
WHERE sortcode=56

UPDATE branch
SET cash = cash + 10000.00
WHERE sortcode=34
COMMIT TRANSACTION
  
```

database management systems  
(DBMS) implements indivisible tasks  
called **transactions**

## The ACID Properties

- **Atomicity** all or nothing
- **Consistency** consistent before → consistent after
- **Isolation** independent of any other transaction
- **Durability** completed transaction are durable

## Transaction Properties: Atomicity

BEGIN TRANSACTION

UPDATE branch

SET cash=cash-10001.00

WHERE sortcode=56

CRASH

<https://tutorcs.com>

## Failure to maintain Atomicity

Suppose that the system crashes half way through processing a cash transfer, and the first part of the transfer has been written to disc

- The database on disc is left in an inconsistent state: the sum of cash should be £137,246.12 but only £127,246.12 recorded
- A DBMS implementing **Atomicity** of transactions would on restart undo the change to branch 56

## Transaction Properties: Consistency

```
BEGIN TRANSACTION
```

```
DELETE FROM branch
```

```
WHERE sortcode=56
```

```
INSERT INTO account
```

```
VALUES (100, 'Smith, J', 'deposit', 5.00, 34)
```

```
END TRANSACTION
```

# Assignment Project Exam Help

<https://tutorcs.com>

### Failure to maintain Consistency

Suppose that a user deletes branch with sortcode 56, and inserts a desposit account number 100 for John Smith at branch sortcode 34

- The database is left in an inconsistent state for two reasons
  - it has three accounts recorded for a branch that appears not to exist, and
  - it has two records for account number 100, with different details for the account
- A DBMS implementing **Consistency** of transactions would forbid both of these changes to the database

## Transaction Properties: Isolation

BEGIN TRANSACTION

UPDATE branch

SET cash=cash-10000.00

WHERE sortcode=56

BEGIN TRANSACTION

```
SELECT SUM(cash) AS net_cash
FROM branch
```

UPDATE branch

SET cash=cash+10000.00

WHERE sortcode=34

END TRANSACTION

END TRANSACTION

Failure to maintain Isolation

Suppose that the system sums the cash in the bank in one transaction, half way through processing a cash transfer in another transaction

- The result of the summation of cash in the bank erroneously reports £127,246.12, whereas the movement of cash always leaves a total of £137,246.12
- A DBMS implementing **Isolation** of transactions ensures that transactions always report results based on the values of committed transactions

## Transaction Properties: Durability

```
BEGIN TRANSACTION
```

```
UPDATE branch
```

```
SET cash=cash -10000.00
```

```
WHERE sortcode=56
```

```
UPDATE branch
```

```
SET cash=cash +10000.00
```

```
WHERE sortcode=34
```

```
END TRANSACTION
```

**CRASH**

Assignment Project Exam Help

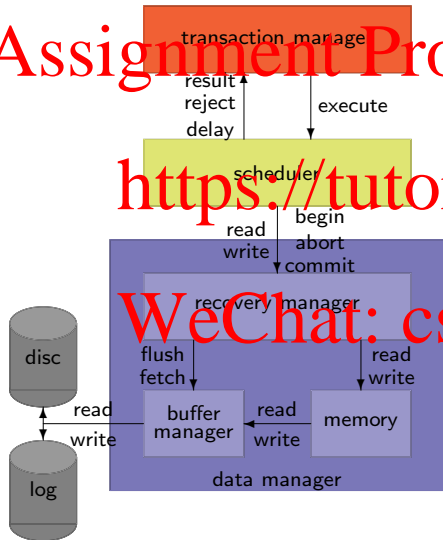
<https://tutorcs.com>

### Failure to maintain Durability

Suppose that the system crashes after informing the user that it has committed the transfer of cash, but has not yet written to disc the update to branch 34

- The database on disc is left in an inconsistent state, with £10,000 'missing'
- A DBMS implementing **Durability** of transactions would on restart complete the change to branch 34 (or alternatively never inform a user of commitment with writing the results to disc).

## DBMS Architecture



## Transaction Manager

BEGIN TRANSACTION

UPDATE branch

SET cash=cash-10000.00

WHERE sortcode=56

UPDATE branch

SET cash=cash+10000.00

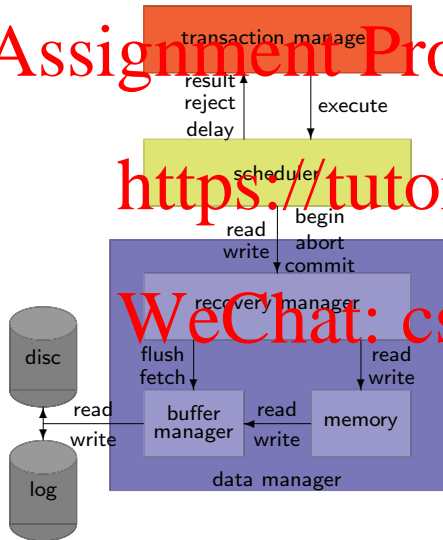
WHERE sortcode=34

END TRANSACTION


 $r_1[b_{56}], w_1[b_{56}], r_1[b_{34}], w_1[b_{34}]$



## DBMS Architecture



## Transaction Manager

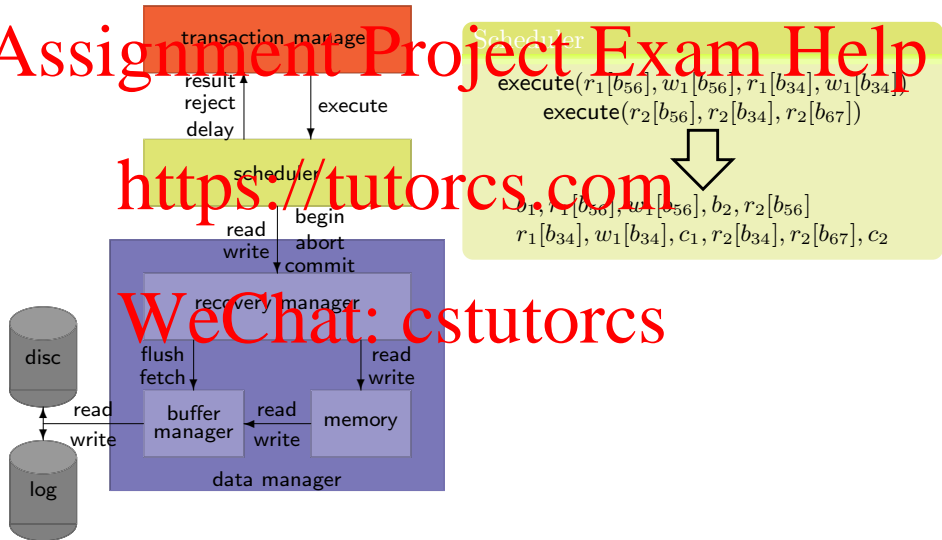
```

BEGIN TRANSACTION
SELECT SUM(cash) AS net_cash
FROM branch
END TRANSACTION

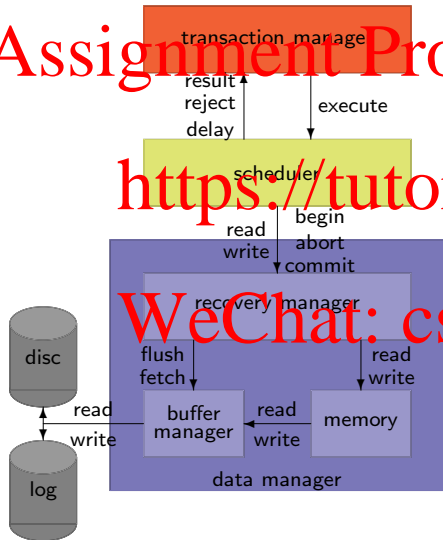
```


 $r_2[b_{56}], r_2[b_{34}], r_2[b_{67}]$

## DBMS Architecture



## DBMS Architecture



## Data Manager

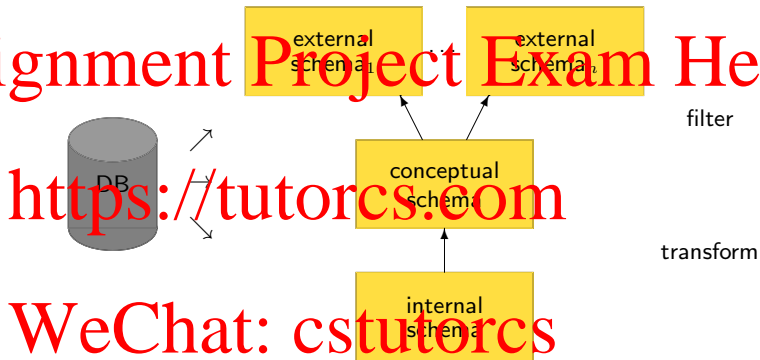
$b_1, r_1[b_{56}], w_1[b_{56}], b_2,$   
 $r_2[b_{56}]$   
 $r_1[b_{34}], w_1[b_{34}],$   
 $c_1, r_2[b_{34}],$   
 $r_2[b_{67}], c_2$



Say  $P_1 = [b_{34}], P_2 = [b_{56}, b_{67}]$

$M_r(P_1), B_r(P_2), D_r(P_2), M_w(P_2), L_w(P_2),$   
 $M_r(P_2),$   
 $M_r(P_1), B_r(P_1), D_r(P_1), M_w(P_1), L_w(P_1),$   
 $L_w(c_1), M_r(P_1),$   
 $M_r(P_2), L_w(c_2), D_w(P_1), D_w(P_2)$

## ANSI/SPARC Model




- ANSI/SPARC model views three levels of abstractions
- **schema** means *structure of the database*

## ANSI/SPARC Model (Internal Schema)

# Assignment Project Exam Help

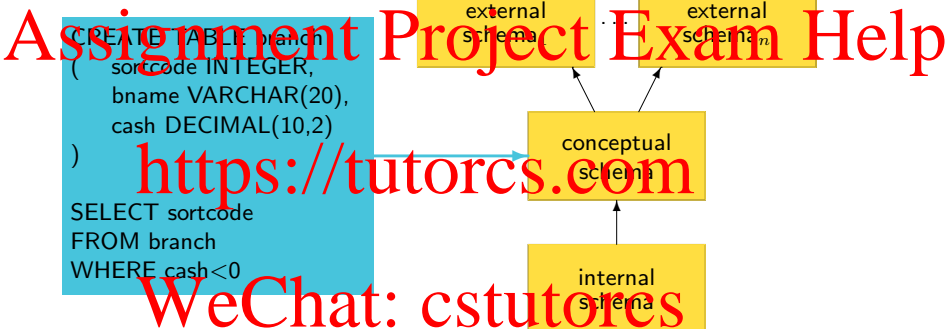
<https://tutorcs.com>

- 2k page size
- B-tree index
- Strings end with char 0



- Describes the physical layout of data


## ANSI/SPARC Model (Conceptual Schema)



- defined in **data definition language (DDL)**
- queried using **data manipulation language (DML)**
- controlled by **database administrator (DBA)**

## ANSI/SPARC Model (External Schema)

```
CREATE VIEW bust
SELECT src code
FROM branch
WHERE cash < 0
```



<https://tutorcs.com>

WeChat: cstutorcs

- Define a schema for a particular user/application

## Course Format

### Schedule

- Three hours combined lectures/tutorials per week, running into week 10
- Coursework that helps you prepare for the exam
- May Exam

### Books

Several good text books on the market. Some that will also cover material in more advanced courses are:

- *Fundamentals of Database Systems*,  
6th Ed, Elmasri and Navathe, Addison Wesley
- *Database Systems: The Complete Book*,  
2nd Ed, Garcia-Molina, Ullman and Widom, Pearson
- *Database Systems*,  
5th Ed, Connolly and Begg, Addison Wesley



## Course Resources

### Course Web Site

<http://www.doc.ic.ac.uk/~ijm/db/>

- Lecture slides
- Example Databases
  - PostgreSQL
  - SQL Server

# Assignment Project Exam Help

# <https://tutorcs.com>

### Resources

- **CATe** course work handout and submission
- **Piazza** discussion forum
- **email** course email list

# WeChat: cstutorcs

If you are not on Level 2 on CATe, nothing works!

## Course Content

### Conceptual Layer: Relational Algebra

- SQL
- Datalog

### Conceptual Layer: Relational Data Model

- Properties of a 'good' schema: keys and normalisation
- Database design using ER models

### Physical Layer: Transaction Processing

- Serialisability
- Recovery and Checkpointing