

Assignment Project Exam Help

Recovery

<https://tutorcs.com>

P.J. McBrien

Imperial College London

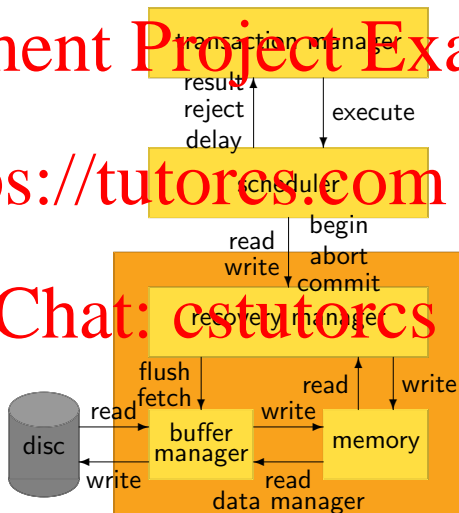
WeChat: cstutorcs

DBMS Architecture

Assignment Project Exam Help

<https://tutor.com>

WeChat: cstutorcs



Recovery Manager (RM)

Assignment Project Exam Help

protect the DBMS against failures

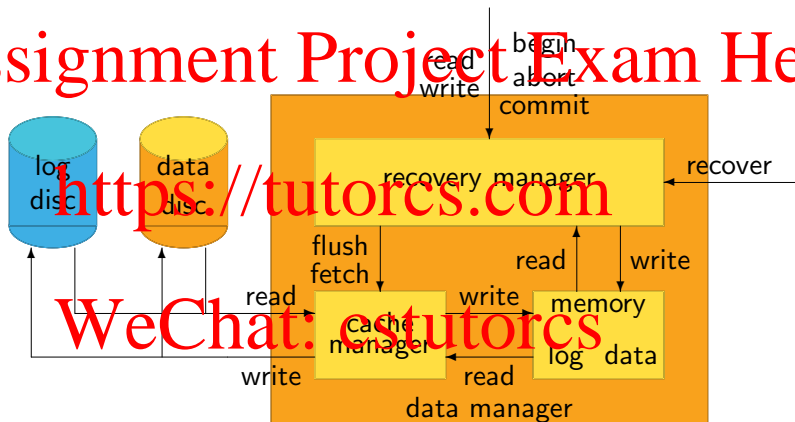
- **system failures** loss of volatile storage

- 1 committed transactions written to disc
 - 2 uncommitted transactions not written to disc
- OR

- 3 sufficient information such that (1) and (2) may be met by a *recovery*

- **media failures** loss of stable storage

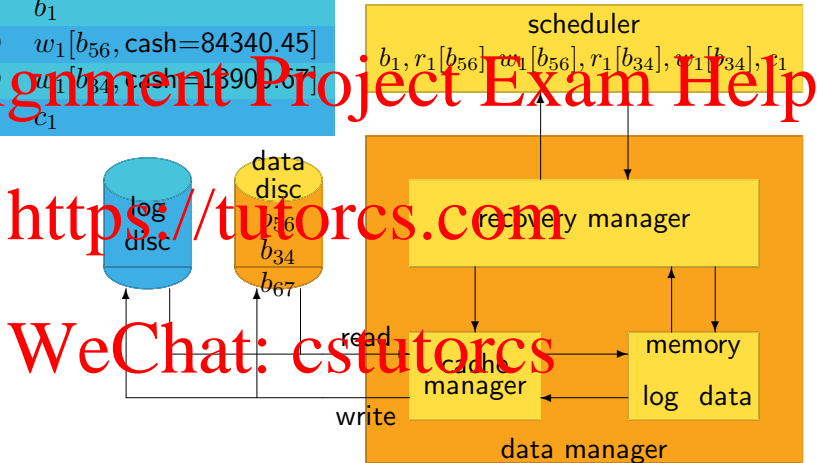
Enhanced Data Manager Architecture



- Need to cache log as well

Need to REDO


LOG	b_1
REDO	$w_1[b_{56}, \text{cash}=84340.45]$
REDO	$w_1[b_{34}, \text{cash}=13900.57]$
LOG	c_1



- REDO required if committed transactions not in stable storage
- must write all REDO to log before commit of transaction

Need to UNDO

LOG	b_1
UNDO	$w_1[b_{56}, \text{cash}=94340.45]$
UNDO	$w_1[b_{34}, \text{cash}=8900.67]$
LOG	c_1



- UNDO required if non-committed transactions in stable storage
- Must flush UNDO to log before corresponding write to data

Quiz 1: Contents of Data Disc After a Transaction

branch		
sortcode	bname	cash
56	'Wimbledon'	94340.45
34	'Goodge St'	18900.67
67	'Strand'	34005.00

```

BEGIN TRANSACTION T1
UPDATE branch
SET cash=cash-10000.00
WHERE sortcode=56

UPDATE branch
SET cash=cash+10000.00
WHERE sortcode=34
COMMIT TRANSACTION T1
  
```

branch ①		
sortcode	bname	cash
56	'Wimbledon'	94340.45
34	'Goodge St'	18900.67
67	'Strand'	34005.00

branch ③		
sortcode	bname	cash
56	'Wimbledon'	84340.45
34	'Goodge St'	18900.67
67	'Strand'	34005.00

branch ②		
sortcode	bname	cash
56	'Wimbledon'	94340.45
34	'Goodge St'	18900.67
67	'Strand'	34005.00

branch ④		
sortcode	bname	cash
56	'Wimbledon'	84340.45
34	'Goodge St'	18900.67
67	'Strand'	34005.00

What must the contents of the branch table on the data disc be after the transaction commits?

A

④

B

① or ④

C

①, ③ or ④

D

①, ②, ③ or ④

Quiz 2: Contents of Log Disc After a Transaction

Data Disc Before Transaction

branch		
sortcode	bname	cash
56	'Wimbledon'	94340.45
34	'Goodge St'	8900.67
67	'Strand'	34005.00

BEGIN TRANSACTION T1

UPDATE branch

SET cash=cash-10000.00

WHERE sortcode=56

UPDATE branch

SET cash=cash+10000.00

WHERE sortcode=34

COMMIT TRANSACTION T1

Data Disc At Commit Time

branch		
sortcode	bname	cash
56	'Wimbledon'	94340.45
34	'Goodge St'	18900.67
67	'Strand'	34005.00

<https://tutorcs.com>

What must be on the log disc after commit time?

A

REDO r_{56} REDO r_{34} UNDO r_{56} UNDO r_{34}

B

REDO r_{56} UNDO r_{34}

C

UNDO r_{34}

D

REDO r_{56}

Before and after images

before image

	branch	
<u>sortcode</u>	bname	cash
56	'Wimbledon'	94340.45
34	'Goodge St'	8900.67
67	'Strand'	34005.00



$w_1[b_{56}]$



	branch	
<u>sortcode</u>	bname	cash
56	'Wimbledon'	84340.45
34	'Goodge St'	8900.67
67	'Strand'	34005.00

after image

Assignment Project Exam Help

<https://tutorcs.com>

WeChat: cstutorcs

- before image allows RM to **undo** $w_1[b_{56}]$
- after image allows RM to **redo** $w_1[b_{56}]$

Database Logs

LOG b_1
 REDO $w_1[b_{56}, \text{cash}=84340.45]$
 REDO $w_1[b_{34}, \text{cash}=18900.67]$
 LOG c_1

LOG b_1
 UNDO $w_1[b_{56}, \text{cash}=94340.45]$
 UNDO $w_1[b_{34}, \text{cash}=8900.67]$
 LOG c_1

<https://tutorcs.com>

LOG b_1
 UNDO $w_1[b_{56}, \text{cash}=94340.45]$
 REDO $w_1[b_{56}, \text{cash}=84340.45]$
 UNDO $w_1[b_{34}, \text{cash}=8900.67]$
 REDO $w_1[b_{34}, \text{cash}=18900.67]$
 LOG c_1

WeChat: cstutorcs

What must a complete REDO/UNDO log contain?

Must contain

- REDO information for each update
- UNDO information for each update
- commit of each transaction

Might contain

- begin of each transaction
 - can be inferred from first REDO/UNDO
 - presence useful to stop search of UNDO records
- abort of each transaction
 - can be inferred from lack of commit
 - presence useful to indicate UNDO already done

Rules for log and data updates

Assignment Project Exam Help

write ahead logging (WAL)

Redo rule

- commit \rightarrow flush log of transaction to disc
- never respond to scheduler before log written

Undo rule:

- flushing uncommitted data \rightarrow flush log of operations

<https://tutorcs.com>

WeChat: cstutorcs

Basic Recovery Procedure

Assignment Project Exam Help



$\Rightarrow w_v[o_1], c_v, w_z[o_2], u_y[o_1], c_y, w_z[o_2] \Rightarrow$



- 1 UNDO \rightarrow Scan back through the log
 - Collect set of committed transactions $C = \{v, y\}$
 - Collect set of incomplete transactions $I = \{x, z\}$
 - Perform UNDO for any transaction in $I = w_z[o_2], w_x[o_2]$
- 2 REDO \rightarrow Scan forward through the log
 - Perform REDO for any transaction in $C = w_v[o_1], w_y[o_1]$



$\Rightarrow \text{UNDO}(w_z[o_2]), \text{UNDO}(w_x[o_2]), \text{REDO}(w_v[o_1]), \text{REDO}(w_y[o_1]) \Rightarrow$



Example of Recovery

Log

LOG	b_4
LOG	w_1
UNDO	$w_1[b_{56}, \text{cash}=94340.45]$
REDO	$w_1[b_{56}, \text{cash}=84340.45]$
LOG	b_2
UNDO	$w_1[b_{34}, \text{cash}=18900.67]$
REDO	$w_2[b_{34}, \text{cash}=8900.67]$
UNDO	$w_2[b_{67}, \text{cash}=34005.00]$
REDO	$w_2[b_{67}, \text{cash}=36005.25]$
LOG	b_7
LOG	c_2
UNDO	$w_1[b_{34}, \text{cash}=8900.67]$
REDO	$w_1[b_{34}, \text{cash}=18900.67]$
UNDO	$w_7[b_{67}, \text{cash}=36005.25]$
REDO	$w_7[b_{67}, \text{cash}=37005.25]$
LOG	c_7
LOG	c_4

Disc Before Recovery

branch			
<u>sortcode</u>	bname	cash	
56	'Wimbledon'	84340.45	
34	'Goodge St'	18900.67	
67	'Strand'	34005.00	

Disc After Recovery

branch			
<u>sortcode</u>	bname	cash	
56	'Wimbledon'	94340.45	
34	'Goodge St'	8900.67	
67	'Strand'	37005.25	

Omitting the REDO Log

Assignment Project Exam Help

If no REDO records kept

must flush committed transactions to data disc

- 1 $C = \emptyset, D = \emptyset$
- 2 Scan the log backwards from the end.
- 3 commit entry \rightarrow add to C
- 4 undo entry for member of $C \rightarrow$ add object to D *without* making changes to the data.
- 5 perform undo entry for object not of member D

Omitting the Undo Log

Assignment Project Exam Help

If no UNDO records kept

transaction must never write uncommitted data

- add **fix** command between RM and CM to stop CM flushing data
- commit is followed by flush or **unfix** of fixed objects

Omitting UNDO and REDO

atomic commit → out of place updating

Quiz 3: Contents of Disc Before Commit if no UNDO log

branch		
sortcode	bname	cash
56	'Wimbledon'	94340.45
34	'Goodge St'	18900.67
67	'Strand'	34005.00

```

BEGIN TRANSACTION T1
UPDATE branch
SET cash=cash-10000.00
WHERE sortcode=56

UPDATE branch
SET cash=cash+10000.00
WHERE sortcode=34
COMMIT TRANSACTION T1
  
```

branch ①		
sortcode	bname	cash
56	'Wimbledon'	94340.45
34	'Goodge St'	8900.67
67	'Strand'	34005.00

branch ③		
sortcode	bname	cash
56	'Wimbledon'	84340.45
34	'Goodge St'	8900.67
67	'Strand'	34005.00

branch ②		
sortcode	bname	cash
56	'Wimbledon'	94340.45
34	'Goodge St'	18900.67
67	'Strand'	34005.00

branch ④		
sortcode	bname	cash
56	'Wimbledon'	84340.45
34	'Goodge St'	18900.67
67	'Strand'	34005.00

What must the contents of the branch table on disc be before the transaction commits?

A

①

B

① or ④

C

④

D

①, ②, ③ or ④

Quiz 4: Contents of Disc After Commit if no REDO log

branch		
sortcode	bname	cash
56	'Wimbledon'	94340.45
34	'Goodge St'	18900.67
67	'Strand'	34005.00

```

BEGIN TRANSACTION T1
UPDATE branch
SET cash=cash-10000.00
WHERE sortcode=56

UPDATE branch
SET cash=cash+10000.00
WHERE sortcode=34
COMMIT TRANSACTION T1
  
```

branch ①		
sortcode	bname	cash
56	'Wimbledon'	94340.45
34	'Goodge St'	8900.67
67	'Strand'	34005.00

branch ③		
sortcode	bname	cash
56	'Wimbledon'	84340.45
34	'Goodge St'	8900.67
67	'Strand'	34005.00

branch ②		
sortcode	bname	cash
56	'Wimbledon'	94340.45
34	'Goodge St'	18900.67
67	'Strand'	34005.00

branch ④		
sortcode	bname	cash
56	'Wimbledon'	84340.45
34	'Goodge St'	18900.67
67	'Strand'	34005.00

What must the contents of the branch table on disc be after the transaction commits?

A

①

B

① or ④

C

④

D

①, ②, ③ or ④

Checkpointing

Assignment Project Exam Help

$\dots, w_x[o_1] \Rightarrow$  $\Rightarrow cp \Rightarrow$  $\Rightarrow w_y[o_1], \dots$

<https://tutorcs.com>

- Forces the database into some known state
- Recovery limited to only look back to checkpoint (or a 'bit' before!)
 - speeds the recovery operation
 - limits the size of log
- The more consistent this known state
 - the easier it is to recover
 - the longer it takes to perform the checkpoint

WeChat: cstutorcs

Commit Consistent Checkpoint

Assignment Project Exam Help

Generating a Commit Consistent Checkpoint

- 1 Stop accepting new transactions
 - 2 Finish existing transactions.
 - 3 Flush all dirty data cache objects to disc.
 - 4 Write a checkpoint to stable log.
- recovery now only needs to scan back to *cp* in log ✓
- possible long hold-up at checkpoint ✗

Cache Consistent Checkpoint

Generating a Cache Consistent Checkpoint

- 1 Suspend all transactions
- 2 Flush all dirty cache objects to disc
- 3 Write a checkpoint + active transactions to stable log

Recovery from Cache Consistent Checkpoint records

- 1 perform UNDOs of non-committed transactions back to *cp*
 - 2 perform UNDO of non-committed transactions before *cp* if they were active at *cp*
 - 3 perform REDOs of committed transactions after *cp*
- could still have delay whilst flushing cached objects

Worksheet: Cache Consistent Checkpoint

```

LOG      b7
UNDO     w7[b67, cash=34005.25]
REDO     w7[b67, cash=36005.25]
LOG      b2
UNDO     w2[b34, cash=10900.67]
REDO     w2[b34, cash=8900.67]
LOG      b6
UNDO     w6[a101, rate=5.25]
REDO     w6[a101, rate=6.00]
LOG      b1
UNDO     w1[b56, cash=94340.45]
REDO     w1[b56, cash=84340.45]
LOG      a7
LOG      cp{1, 2, 6}
      ⋮

```

```

      ⋮
UNDO     w6[a119, rate=5.50]
REDO     w6[a119, rate=6.00]
LOG      c6
UNDO     w2[b67, cash=34005.00]
REDO     w2[b67, cash=36005.25]
LOG      b8
LOG      c2
UNDO     w1[b34, cash=8900.67]
REDO     w1[b34, cash=18900.67]
LOG      b9
UNDO     w9[b67, cash=36005.00]
REDO     w9[b67, cash=20000.00]
LOG      c9

```

Fuzzy Checkpointing

Generating a Fuzzy Checkpoint


- 1 Suspend all transactions
- 2 Flush any dirty cache objects to disc not flushed in previous *cp*
- 3 Write a checkpoint + active transactions to stable log

Recovery from Fuzzy Checkpoint records

Recovery works like cache consistent checkpoint, but working with penultimate *cp*


- 1 perform UNDOs of non-committed transactions back to penultimate *cp*
- 2 perform UNDO of non-committed transactions before penultimate *cp* if they were active at *cp*
- 3 perform REDOs of committed transactions after penultimate *cp*

Media Failures: Mirroring (RAID-1)



- Keep more than one active copy of data and log
- Writes sent to both
- Read from either

Media Failures: Dumping



- 'tape' might also be a external file server, removable HD, *etc.*
- To use normal OS backup procedure
 - DBMS must not be still running
 - raw partition must not be used

Checkpoints and Dumps


Assignment Project Exam Help

- Dump must do a checkpoint
- Restore involves:
 - 1 copy tape to disc
 - 2 undo transactions active at the archive time
 - 3 redo transactions that committed after the archive
- commit consistent checkpoint obvious choice

<https://tutorcs.com>

WeChat: cstutorcs

Media Failures: Archive Database



- mirror log, but only have one active database
- periodically archive updates onto archive database
- failure of active database disc involves restore of archive database using logs

THE END

Assignment Project Exam Help

- Content of the course is what has been presented in the lectures
- Revise by reviewing worksheets and courseworks
- 2011 exam papers onwards set to current syllabus
- Older exam questions mostly apply, but there is more emphasis on RA and SQL, less on concurrency.

<https://tutorcs.com>
WeChat: cstutorcs